

Maxence PACHOT

Data Engineer · Platform & DataOps

Azure · Databricks · Terraform · dbt

✉ pachot.maxence@gmail.com

☎ +41 76 391 15 69

🌐 pachotmaxence.com

🌐 maxence-pachot

🐙 MaxencePachot

Renens (VD), Switzerland · Swiss B residence permit · Category B driving licence



PROFILE

Three and a half years in Swiss IT consulting, across pharmaceuticals, social insurance and healthcare. I design and maintain Azure data platforms, from the Terraform infrastructure to the Databricks lakehouse, CI/CD included. I am currently the technical lead on a pharmaceutical data platform.

PROFESSIONAL EXPERIENCE

Engineer, Business Line Data Analytics & AI · ELCA Informatique SA

03.2023 – present

Swiss IT consulting firm · Lausanne, Switzerland · data platform and database projects for customers in pharmaceuticals, social insurance and healthcare

Azure lakehouse for a pharmaceutical group

02.2026 – present

Data Platform Engineer, technical lead since 09.2026 (2 developers)

- The group went from an empty Azure subscription to a production lakehouse feeding Power BI. I built its four layers: infrastructure, ingestion, ETL and CI/CD.
- Azure landing zone in Terraform: Hub & Spoke, private by default, with the Azure DevOps pipelines that deploy it on agents inside the network and OIDC federation rather than stored secrets.
- Ingestion from an on-premises SQL Server through Azure Data Factory, into a medallion lakehouse on Databricks and Unity Catalog: SCD1/SCD2 merge engine and conformed dimensions shared across domains. Finance and marketing in production.
- As technical lead, I review the code of the shared foundation, on which two developers each deliver their own business domain, and I run the workshops with the business. I moved the logic out of the notebooks into a tested Python package, and made review mandatory on the shared core.

Terraform · Azure · Databricks · Unity Catalog · PySpark · Delta Lake · Azure Data Factory · SQL Server · Azure DevOps · pytest

Azure data platform for a healthcare group

06.2026 – present

Platform Engineer / DataOps

- I deployed the group's Terraform data platform for a healthcare customer, on a brand-new Azure subscription: two-VNet network bootstrap, private CI runners, Terraform state behind a private endpoint.
- The platform could only deploy to a single region and a single tenant. I built multi-region support into the shared engine, and the cross-tenant bridge the split of responsibilities requires: we run the compute, the customer keeps its own storage.
- The Unity Catalog governance layer, then a pilot dbt pipeline on a Databricks SQL warehouse, packaged as a Databricks Asset Bundle and run by Azure DevOps pipelines I built. The platform runs in production.

Terraform · Azure · Databricks · Unity Catalog · dbt · Databricks Asset Bundles · Azure DevOps · Bicep · Bash

Migrating a social insurance data platform to dbt

12.2023 – present

Data Engineer

- I built part of the legacy chain, SSIS packages and stored procedures, then migrated it entirely to dbt on SQL Server with one colleague. It now runs in production.
- That layer now lives inside a shared dbt library of over a thousand models, released as a versioned package that each compensation fund imports, instead of an ETL rebuilt per fund.
- I own 5 of its 26 business domains, contributions and tax reporting among them, from the DB2 extract through to the export, with regression and integrity tests on each.
- Data models for the staging and target layers, sqlfluff linting on changed files, and the business requirements gathered with the stakeholders. The whole thing is orchestrated with Dagster.

Infrastructure cost reduction on a healthcare platform

10.2023 – 01.2024

Data Engineer / DevOps

- Migrated the workloads from Kubernetes to Azure serverless containers, dividing the infrastructure bill by roughly ten.
- Containerised the services with Docker, automated their deployment through Azure DevOps, and added a new data source with the Python and FastAPI features consuming it.

Azure serverless containers · Docker · Python · FastAPI · MongoDB · Azure DevOps

Oracle database administration

03.2023 – present

Database Administrator

- Deploy and operate Oracle databases from 11g to 23c across customer environments, handle incidents through JIRA, and write and tune PL/SQL.
- Automated fleet-wide tablespace and disk monitoring in Python, so saturation is caught before it becomes an incident.

Oracle 11g–23c · PL/SQL · SQL · Python · Linux · JIRA

Full Stack Developer, internship · SOFI Groupe

04.2022 – 08.2022

France

- Built a web interface and API letting external partners manage buyback catalogues, and a system reconciling discrepancies between two company databases, integrated with the existing ETL processes.

PHP · SQL Server · MySQL · Python · Git

SKILLS, BY HOW OFTEN I USE THEM

Every day

Azure: Databricks, Data Factory, ADLS Gen2, Key Vault, VNet and private endpoints, NAT gateway, Log Analytics, Entra ID service principals · **Terraform:** reusable modules, remote state, multi-environment layouts, tfint, terraform-docs, pre-commit · **Databricks:** Unity Catalog, Delta Lake, Delta Live Tables and Lakeflow, Auto Loader, Asset Bundles, SQL Warehouses, medallion architecture, SCD1/SCD2 merge · **dbt** · **CI/CD:** Azure DevOps, GitHub Actions, self-hosted runners in private networks, OIDC and workload identity federation, branch policies · **Python** (PySpark, pytest, ruff), **SQL**, T-SQL, PL/SQL, Bash · AI-assisted development with Claude Code

Solid

Oracle 11g–23c administration: deployment, performance tuning, PL/SQL, monitoring · **SQL Server** and dbt on SQL Server · Bicep · Linux

Working knowledge

Dagster · Docker · FastAPI · MongoDB · React · Vite

CERTIFICATIONS

Databricks Certified Data Engineer Associate	Nov 2024
Microsoft Certified: Azure Fundamentals (AZ-900)	Sept 2023

EDUCATION

Master 2 MIAGE / MBDS · Université Côte d'Azur	2022 – 2023
Engineering degree · ESTIA, Bidart (FR)	2020 – 2023
CPGE PCSI / PC · Lycée Jacques Amyot, Melun (FR)	2018 – 2020

LANGUAGES

French: native · English: C1 (TOEIC 910)